

Akshita Bhamidimarri

Analytics Platform on Azure Databricks

Capstone project - 2

1) Creating a storage account

The screenshot shows the Microsoft Azure portal interface. The browser tabs include 'raw - Microsoft', 'aksssstore - Mic', 'Catalog Explore', 'de-with-azure', 'Project 02 - Go', and 'capstone-2 - Go'. The address bar shows the URL: <https://portal.azure.com/#@michaelkloudkrafthotmail.onmicrosoft.com/resource/subscriptions/d008c7cc-9795-4292-bf79-74ae52cda63d/r...>. The page title is 'aksssstore' and it is identified as a 'Storage account'. The left sidebar shows the 'Overview' tab selected, with options like Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, Storage browser, Partner solutions, Resource visualizer, Data storage, Containers, File shares, Queues, and Tables. The main content area shows the 'Essentials' section with a table of account details:

Property	Value
Resource group	LabsKraft2027
Location	eastus
Subscription	LabsKraft
Subscription ID	d008c7cc-9795-4292-bf79-74ae52cda63d
Disk state	Available
Tags	(edit)
Performance	Standard
Replication	Locally-redundant storage (LRS)
Account kind	StorageV2 (general purpose v2)
Provisioning state	Succeeded
Created	8/26/2025, 10:00:16 AM

Below the Essentials section, there are tabs for Properties, Monitoring, Capabilities (5), Recommendations (1), Tutorials, and Tools + SDKs. The 'Data Lake Storage' section shows 'Hierarchical namespace' as 'Enabled'. The 'Security' section shows 'Require secure transfer for REST API' as 'Enabled'. The bottom of the page shows a weather widget for 'Rainy days ahead 24°C' and the system clock for '10:17 AM 8/26/2025'.


2) Creating a container inside the storage account

The screenshot shows the Microsoft Azure portal interface for the 'akscontainer' container. The browser tabs include 'raw - Microsoft', 'akscontainer - b', 'Catalog Explore', 'de-with-azure', 'Project 02 - Go', and 'capstone-2 - Go'. The address bar shows the URL: https://portal.azure.com/#view/Microsoft_Azure_Storage/ContainerMenuBlade/~/overview/storageAccountid/%2Fsubscriptions%2Fd008c7.... The page title is 'akscontainer' and it is identified as a 'Container'. The left sidebar shows the 'Overview' tab selected, with options like Diagnose and solve problems, Access Control (IAM), and Settings. The main content area shows the 'Overview' section with the 'Authentication method' set to 'Access key (Switch to Microsoft Entra user account)'. Below this, there is a search bar and a dropdown menu for 'Only show active objects'. A table shows the contents of the container:

Name	Last modified	Access tier	Blob type	Size	Lease state
products	8/26/2025, 10:13:48 AM				...
sales	8/26/2025, 10:14:02 AM				...


The bottom of the page shows a weather widget for 'Rainy days ahead 24°C' and the system clock for '10:18 AM 8/26/2025'.

3) Add files to two folders called sales and products



The screenshot shows the Microsoft Azure portal interface. The breadcrumb navigation is 'Home > aksstore | Containers > akscontainer'. The left sidebar shows 'Overview', 'Diagnose and solve problems', 'Access Control (IAM)', and 'Settings'. The main content area shows the 'akscontainer' container with the 'products' folder selected. The authentication method is 'Access key'. A search bar is present with the text 'Search blobs by prefix (case-sensitive)'. Below the search bar, it says 'Showing all 2 items'. A table lists the contents of the 'products' folder:

	Name	Last modified	Access tier	Blob type	Size	Lease state
☐	[.]					...
☐	products.csv	8/26/2025, 10:14:59 AM	Hot (Inferred)	Block blob	36.69 KiB	Available
☐	products.json	8/26/2025, 10:15:00 AM	Hot (Inferred)	Block blob	138.2 KiB	Available



The screenshot shows the Microsoft Azure portal interface. The breadcrumb navigation is 'Home > aksstore | Containers > akscontainer'. The left sidebar shows 'Overview', 'Diagnose and solve problems', 'Access Control (IAM)', and 'Settings'. The main content area shows the 'akscontainer' container with the 'sales' folder selected. The authentication method is 'Access key'. A search bar is present with the text 'Search blobs by prefix (case-sensitive)'. Below the search bar, it says 'Showing all 5 items'. A table lists the contents of the 'sales' folder:

	Name	Last modified	Access tier	Blob type	Size	Lease state
☐	[.]					...
☐	sales_day_1.csv	8/26/2025, 10:15:24 AM	Hot (Inferred)	Block blob	50.54 KiB	Available
☐	sales_day_2.csv	8/26/2025, 10:15:24 AM	Hot (Inferred)	Block blob	50.48 KiB	Available
☐	sales_day_3.csv	8/26/2025, 10:15:24 AM	Hot (Inferred)	Block blob	50.46 KiB	Available
☐	sales_day_4.csv	8/26/2025, 10:15:23 AM	Hot (Inferred)	Block blob	50.5 KiB	Available
☐	sales_day_5.csv	8/26/2025, 10:15:23 AM	Hot (Inferred)	Block blob	50.51 KiB	Available

4) Create an azure connector for databricks

The screenshot shows the Microsoft Azure portal interface. The browser address bar displays the URL: <https://portal.azure.com/#@michaelkloudkraft@hotmail.onmicrosoft.com/resource/subscriptions/d008c7cc-9795-4292-bf79-74ae52cda63d/r...>. The page title is "aksconnect" with the subtitle "Access Connector for Azure Databricks". The left sidebar contains a navigation menu with options: Overview, Activity log, Access control (IAM), Tags, Resource visualizer, Settings, Automation, and Help. The main content area shows the "Overview" tab for the "aksconnect" resource. It includes a search bar, "Refresh" and "Delete" buttons, and a "JSON View" link. The "Essentials" section displays the following information:

- Resource group: [LabsKraft2027](#) (move)
- Location: East US
- Subscription: [LabsKraft](#) (move)
- Subscription ID: d008c7cc-9795-4292-bf79-74ae52cda63d
- Tags: [\(edit\)](#) [Add tags](#)
- State: Succeeded
- Resource ID: /subscriptions/d008c7cc-9795-4292-bf79-74ae52cda63d/resourceGroups/...

The bottom of the screen shows a Windows taskbar with a search bar, task icons, and a system tray displaying the time as 10:21 AM on 8/26/2025.

5) Give the azure connector Storage Blob Data Contributor access

The screenshot shows the Microsoft Azure portal interface for the "aksstore" storage account. The browser address bar displays the URL: <https://portal.azure.com/#@michaelkloudkraft@hotmail.onmicrosoft.com/resource/subscriptions/d008c7cc-9795-4292-bf79-74ae52cda63d/r...>. The page title is "aksstore | Access Control (IAM)" with the subtitle "Storage account". The left sidebar contains a navigation menu with options: Overview, Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, Storage browser, Partner solutions, Resource visualizer, Data storage, Containers, File shares, Queues, and Tables. The main content area shows the "Access Control (IAM)" tab for the "aksstore" storage account. It includes a search bar, "Add", "Download role assignments", "Edit columns", "Refresh", "Delete", and "Feedback" buttons. The "Contributor (1)" section shows the following role assignments:

Role	Principal	Role	Source
Owner	Tredence LabsKraft8 a5f77ccd-d0a3-499...	Owner	Resource group (Inher... View)
Owner	Tredence LabsKraft9 e16a3d91-71e7-461...	Owner	Resource group (Inher... View)

The "Monitoring Contributor (1)" section shows the following role assignment:

Role	Principal	Role	Source
Monitoring Contributor	SherlockShinde f837afc-8713-4ff0...	Monitoring Contributor	Subscription (Inherited) None

The "Storage Blob Data Contributor (1)" section shows the following role assignment:

Role	Principal	Role	Source
Storage Blob Data Contributor	aksconnect b29373db-2972-42...	Storage Blob Data Contrib...	This resource Add

The "User Access Administrator (1)" section shows the following role assignment:

Role	Principal	Role	Source
User Access Administrator	Michael A e1fb5008-9e9f-4e28...	User Access Administrator	Root (Inherited) None

The bottom of the screen shows a Windows taskbar with a search bar, task icons, and a system tray displaying the time as 10:24 AM on 8/26/2025.

6) Create a new external location for the new catalog

The screenshot displays the Databricks Catalog Explorer interface for an external location named 'aksloc'. The interface is divided into a left sidebar with navigation options (Workspace, Recents, Catalog, Jobs & Pipelines, Compute, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie, Alerts, Query History, SQL Warehouses, Data Engineering, Job Runs, Data Ingestion, AI/ML) and a main content area. The main content area shows the 'Overview' tab for 'aksloc', which includes a description, a table of credentials, and a 'Test connection' button.

Credential Table:

Credential	Value
aksloc_azuremanagedidentity_1756184085081	
URL	abfss://akscontainer@akssstore.dfs.core.windows.net/
Limit to read-only use	Disabled

About this external location:


- Owner: Tredence LabsKraft2
- Fallback mode: ☐

The 'Test connection' dialog box is open, showing the following results:


- Location Type: Directory
- Success - Read
- Success - List
- Success - Write
- Success - Delete
- Success - Path Exists
- Success - Hierarchical Namespace Enabled
- All Permissions Confirmed: The associated Storage Credential grants permission to perform all necessary operations.

The 'Done' button is visible at the bottom right of the dialog box.

7) Create the new catalog using the newly created external location



8) Grant users specific access to specific roles



aksstore - Mic X aksconnect - Mi X aks_proj2 - Dat X Users - Settings X de-with-azure- X Project 02 - Go X capstone-2 - Go X

https://adb-2262725646607275.15.azuredatabricks.net/settings/workspace/identity-and-access/users?o=2262725646607275

Microsoft Azure databricks Search data, notebooks, recents, and more... CTRL + P tred_d1

+ New

- Workspace
- Recents
- Catalog
- Jobs & Pipelines
- Compute
- Marketplace
- SQL
- SQL Editor
- Queries
- Dashboards
- Genie
- Alerts
- Query History
- SQL Warehouses
- Data Engineering
- Job Runs
- Data Ingestion
- AI/ML

Settings

- Workspace admin
- Appearance
- Identity and access
- Security
- Compute
- Development
- Notifications
- Advanced
- User
- Profile
- Preferences
- Developer
- Linked accounts
- Notifications

Workspace settings > Identity and access > Users

Filter users 9 total Add user

Add user

Enter the email address of your new user.

New user email

akshitab2609@gmail.com

Back Add

✓	tredence12@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft12
✓	tredence1@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft1
✓	tredence2@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft2
✓	tredence3@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft3
✓	tredence4@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft4

24°C Mostly cloudy Search 10:35 AM 8/26/2025

aksstore - Mic X aksconnect - Mi X aks_proj2 - Dat X Users - Settings X de-with-azure- X Project 02 - Go X capstone-2 - Go X

https://adb-2262725646607275.15.azuredatabricks.net/settings/workspace/identity-and-access/users?o=2262725646607275

Microsoft Azure databricks Search data, notebooks, recents, and more... CTRL + P tred_d1

+ New

- Workspace
- Recents
- Catalog
- Jobs & Pipelines
- Compute
- Marketplace
- SQL
- SQL Editor
- Queries
- Dashboards
- Genie
- Alerts
- Query History
- SQL Warehouses
- Data Engineering
- Job Runs
- Data Ingestion
- AI/ML

Settings

- Workspace admin
- Appearance
- Identity and access
- Security
- Compute
- Development
- Notifications
- Advanced
- User
- Profile
- Preferences
- Developer
- Linked accounts
- Notifications

Workspace settings > Identity and access > Users

Filter users 12 total Add user

Status	Email	Name
✓	akshi.bham@gmail.com	
✓	akshitab.2603@gmail.com	
✓	akshitab2609@gmail.com	
✓	labskraft77@michaelkloudkrafthotmail.onmicrosoft.com	labskraft77
✓	michael.kloudkraft@hotmail.com	Michael A
✓	michael.kloudkraft_hotmail.com#ext@michaelkloudkrafthotmail.onmicrosoft.com	Michael A
✓	rohit@michaelkloudkrafthotmail.onmicrosoft.com	RohitY
✓	tredence12@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft12
✓	tredence1@michaelkloudkrafthotmail.onmicrosoft.com	Tredence LabsKraft1

Top Stories Exporters race a...

Search 10:37 AM 8/26/2025

Grant on aks_proj2

Granted privileges will be inherited by applicable objects (e.g. schemas, tables) in this catalog. [Learn more](#)

Principals

akshi.bham@gmail.com X

Privilege presets

Custom

Prerequisite	Read	Create
☒ USE CATALOG	☒ EXECUTE	☒ CREATE FUNCTION
☒ USE SCHEMA	☒ READ VOLUME	☒ CREATE MATERIALIZED VIEW
	☒ SELECT	☒ CREATE MODEL
Metadata	Edit	☒ CREATE MODEL VERSION
☒ APPLY TAG	☒ MODIFY	☒ CREATE SCHEMA
☒ BROWSE	☒ REFRESH	☒ CREATE TABLE
	☒ WRITE VOLUME	☒ CREATE VOLUME

Cancel Confirm

Grant on aks_proj2

Granted privileges will be inherited by applicable objects (e.g. schemas, tables) in this catalog. [Learn more](#)

Principals


akshitab.2603@gmail.com X

Privilege presets

Custom


Prerequisite	Read	Create
☒ USE CATALOG	☒ EXECUTE	☐ CREATE FUNCTION
☒ USE SCHEMA	☒ READ VOLUME	☐ CREATE MATERIALIZED VIEW
	☒ SELECT	☐ CREATE MODEL
Metadata	Edit	☐ CREATE MODEL VERSION
☐ APPLY TAG	☐ MODIFY	☒ CREATE SCHEMA
☐ BROWSE	☒ REFRESH	☒ CREATE TABLE
	☐ WRITE VOLUME	☒ CREATE VOLUME


Cancel Confirm



8) Create schema in the catalog and then start the pipeline

```
%sql
USE CATALOG aks_proj2;
create schema if not exists aks_proj2.sales_analytics;
```





```
import dlt
from pyspark.sql.functions import col, to_date, sum
from pyspark.sql.types import StructType, StructField, StringType,
IntegerType, DoubleType

# -----
# DEFINE SCHEMAS
# -----

# SALES SCHEMA
sales_schema = StructType([
    StructField("sale_id", StringType(), True),
    StructField("product_id", StringType(), True),
    StructField("region", StringType(), True),
    StructField("store_id", StringType(), True),
    StructField("quantity", IntegerType(), True),
    StructField("sales_amount", DoubleType(), True),
    StructField("sale_date", StringType(), True),
    StructField("supplier_last", StringType(), True) # will be cast later
])

# PRODUCTS SCHEMA (Step 2: Added supplier_name)
product_schema = StructType([
    StructField("product_id", StringType(), True),
    StructField("product_name", StringType(), True),
    StructField("category", StringType(), True),
    StructField("price", DoubleType(), True),
    StructField("discount_rate", DoubleType(), True),
    StructField("supplier_name", StringType(), True) # ✅ New column added
])

# -----
```

```

# BRONZE LAYER (Raw Ingest)
# -----

@dlt.table(name="bronze_sales")
def bronze_sales():
    return (
        spark.readStream.format("cloudFiles")
            .option("cloudFiles.format", "csv")
            .option("cloudFiles.schemaLocation",
"abfss://akscontainer@aksssstore.dfs.core.windows.net/schemas/sales")
            .schema(sales_schema)
            .load("abfss://akscontainer@aksssstore.dfs.core.windows.net/sales/")
    )

@dlt.table(name="bronze_products")
def bronze_products():
    return (
        spark.readStream.format("cloudFiles")
            .option("cloudFiles.format", "json")
            .option("multiLine", "true")
            .option("cloudFiles.schemaLocation",
"abfss://akscontainer@aksssstore.dfs.core.windows.net/schemas/products")
            .schema(product_schema)

.load("abfss://akscontainer@aksssstore.dfs.core.windows.net/products/")
    )

# -----
# SILVER LAYER (Cleansed + Validated)
# -----

@dlt.table(name="silver_sales")
@dlt.expect("Valid sales amount", "sales_amount >= 0")
@dlt.expect("Valid quantity", "quantity > 0")
def silver_sales():
    df = dlt.read("bronze_sales")
    return (
        df.dropna()
            .withColumn("sales_amount", col("sales_amount").cast("double"))
            .withColumn("quantity", col("quantity").cast("int"))
            .withColumn("sale_date", to_date("sale_date", "yyyy-MM-dd"))
    )

```

```

    )

@dlt.table(name="silver_products")
@dlt.expect("Valid price", "price > 0")
@dlt.expect("Valid discount rate", "discount_rate >= 0 AND discount_rate <=
1")
def silver_products():
    df = dlt.read("bronze_products")
    return (
        df.dropna(subset=["product_id", "product_name", "category", "price",
"discount_rate"])
        .dropDuplicates(["product_id"])
    )

# -----
# GOLD LAYER (Aggregated Business View)
# -----

@dlt.table(name="gold_sales_summary")
def gold_sales_summary():
    sales = dlt.read("silver_sales").withColumn("product_id",
col("product_id").cast("string"))
    products = dlt.read("silver_products")

    joined_df = sales.join(products, on="product_id", how="inner")

    return (
        joined_df
        .withColumn("effective_price", col("price") * (1 -
col("discount_rate")))
        .withColumn("discount_amount", col("price") * col("discount_rate") *
col("quantity"))
        .withColumn("total_effective_price", col("effective_price") *
col("quantity"))
        .groupBy("sale_date", "region", "product_id", "product_name",
"category")
        .agg(
            sum("quantity").alias("total_units_sold"),
            sum("sales_amount").alias("total_revenue"),
            sum("discount_amount").alias("total_discount_given"),
            sum("total_effective_price").alias("total_effective_price")

```

9) Make change to the schema and try running the pipeline, it will show an error that schema has been changed

Microsoft Azure databricks

Jobs & Pipelines > aks_proj2 ☆ Send feedback

Run: 8/26/2025, 4:26:29 PM · Canceled This update has already finished Show the latest update

Graph List

Pipeline details Update details

Pipeline ID cffbcb4e-8bf5-46c6-a94a-34a8c9524431

Pipeline type ETL pipeline

Source code /Users/tredence2@michaelkloudkraft@hotmail.onmicrosoft.com/aks2

Event log Query history


All Info Warning Error Filter...

1 minute ago flow_progress Flow 'aks_proj2.sales_analytics.products_clean' is STARTING.

1 minute ago flow_progress Flow 'aks_proj2.sales_analytics.products_clean' is RUNNING.

51 seconds ago flow_progress Flow 'aks_proj2.sales_analytics.products_clean' has COMPLETED.

49 seconds ago update_progress Update 28155c has been cancelled due to a schema change in aks_proj2.sales_analytics.sales_raw, and will be restar...



10) Put the clean code and try running the a new pipeline, it will work

```
import dlt
```

```

from pyspark.sql.functions import *
from pyspark.sql.types import *

# Updated schema locations
sales_schema_location =
"abfss://akscontainer@aksssstore.dfs.core.windows.net/schema/sales/"
products_csv_schema_location =
"abfss://akscontainer@aksssstore.dfs.core.windows.net/schema/products_csv/"
products_json_schema_location =
"abfss://akscontainer@aksssstore.dfs.core.windows.net/schema/products_json/"

# Ingest raw sales data
@dlt.table(
    name="proj2_sales_raw_data",
    comment="Raw sales data ingested from CSV files with automated schema
inference",
    table_properties={"quality": "bronze"}
)
def proj3_sales_raw_data():
    return (
        spark.readStream.format("cloudFiles")
        .option("cloudFiles.format", "csv")
        .option("header", "true")
        .option("cloudFiles.schemaLocation", sales_schema_location)
        .option("cloudFiles.schemaEvolutionMode", "addNewColumns")
        .load("abfss://akscontainer@aksssstore.dfs.core.windows.net/sales/")
    )

# Products CSV ingestion
@dlt.table(
    name="proj2_products_csv_raw",
    comment="Raw product data from CSV with automated schema inference",
    table_properties={"quality": "bronze"}
)
def proj3_products_csv_raw():
    return (
        spark.readStream.format("cloudFiles")
        .option("cloudFiles.format", "csv")
        .option("header", "true")
        .option("cloudFiles.schemaLocation", products_csv_schema_location)
        .option("cloudFiles.schemaEvolutionMode", "addNewColumns")

```

```

.load("abfss://akscontainer@aksssstore.dfs.core.windows.net/products/")
)

# Products JSON ingestion
@dlt.table(
    name="proj3_products_json_raw",
    comment="Raw product data from JSON with automated schema inference",
    table_properties={"quality": "bronze"}
)

def proj3_products_json_raw():
    return (
        spark.readStream.format("cloudFiles")
        .option("cloudFiles.format", "json")
        .option("multiLine", "true")
        .option("cloudFiles.schemaLocation", products_json_schema_location)
        .option("cloudFiles.schemaEvolutionMode", "addNewColumns")

.load("abfss://akscontainer@aksssstore.dfs.core.windows.net/products/")
)

# Unified view of products
@dlt.view(name="proj2_products_combined_raw")
def proj3_products_combined_raw():
    return dlt.read_stream("proj3_products_csv_raw").unionByName(
        dlt.read_stream("proj3_products_json_raw"), allowMissingColumns=True
    )

# Cleaned sales data
@dlt.table(
    name="proj2_sales_cleaned",
    comment="Cleaned sales data with data quality checks",
    table_properties={"quality": "silver"}
)

@dlt.expect_or_drop("valid_sales_amount", "sales_amount >= 0")
@dlt.expect_or_drop("valid_quantity", "quantity > 0")
def proj3_sales_cleaned():
    return (
        dlt.read_stream("proj3_sales_raw_data")
        .withColumn("sale_date", to_date(col("sale_date"), "yyyy-MM-dd"))
        .na.drop()

```

```

    )

# Cleaned products data
@dlt.table(
    name="proj2_products_cleaned",
    comment="Cleaned products data with schema evolution support",
    table_properties={"quality": "silver"}
)

def proj3_products_cleaned():
    return (
        dlt.read_stream("proj2_products_combined_raw")
        .withColumn("discount_rate", coalesce(col("discount_rate"),
lit(0.0)))
    )

# Final sales summary
@dlt.table(
    name="proj2_sales_kpi_summary",
    comment="Daily sales summary by product and region",
    table_properties={"quality": "gold"}
)

def proj3_sales_kpi_summary():
    sales = dlt.read_stream("proj2_sales_cleaned")
    products = dlt.read("proj2_products_cleaned")
    return (
        sales.join(products, "product_id", "left")
        .groupBy("region", "sale_date", "product_id", "product_name",
"category")
        .agg(
            sum("sales_amount").alias("daily_revenue"),
            sum("quantity").alias("units_sold"),
            avg("price").alias("avg_price"),
            avg("discount_rate").alias("avg_discount")
        )
    )

```

Microsoft Azure databricks

Search data, notebooks, recents, and more... CTRL + P

+ New

Workspace

Recents

Catalog

Jobs & Pipelines

Compute

Marketplace

SQL

SQL Editor

Queries

Dashboards

Genie

Alerts

Query History

SQL Warehouses

Data Engineering

Job Runs

Data Ingestion

AI/ML

aks2

aks2_new

Untitled Notebook 2025-08-26 17:22:09

File Edit View Run Help Python Tabs: ON

Last edit was 1 minute ago

Validate Start aks_proj2new Share

```
100 products = sql_ctx.read('proj2_products_cleaned',
101
102 return (
103     sales.join(products, "product_id", "left")
104     .groupBy("region", "sale_date", "product_id", "product_name", "category")
105     .agg(
106         sum("sales amount").alias("daily_revenue")
107     )
108 )
```

DLT graph DLT event log DLT query history

Streaming table

proj2_products_cs...
Completed - 7s
8K 0

Streaming table

proj2_sales_raw_d...
Completed - 7s
5K 0

Streaming table

proj3_products_js...
Completed - 2s
- -

View

proj2_products_co...

Streaming table

proj2_products_cle...
Completed - 6s
9K 0

Streaming table

proj2_sales_cleaned
Completed - 6s
0 0


Streaming table

proj2_sales_kpi_su...
Completed - 43s
0 0

26°C Partly sunny

Search

5:36 PM 8/26/2025



11) Connect with power BI to create dashboards

The image is a composite screenshot showing the process of connecting Power BI to Azure Databricks. The top portion displays the Power BI Desktop interface with the 'Azure Databricks' connection dialog box open. The dialog contains the following fields: 'Server Hostname' (adb-2262725646607275.15.azuredatabricks.net), 'HTTP Path' (/sql/1.0/warehouses/624b8981d8e6e924), 'Advanced Options (optional)' section with 'Default catalog (optional)' (Example: abc), 'Database (optional)' (Example: abc), 'Automatic Proxy Discovery (optional)' (dropdown), and 'Native query (Requires: Default catalog) (optional)' (Example: select * from db.schemaname.tablename). The bottom portion of the screenshot shows the Databricks web interface in a browser. The page title is 'Serverless Starter Warehouse'. The 'Connection details' tab is active, showing the same 'Server hostname' and 'HTTP path' as the Power BI dialog. It also displays the 'JDBC URL' (2.6.25 or later) and a 'JDBC connection string' (jdbc:databricks://adb-2262725646607275.15.azuredatabricks.net:443/default;transportMode=http;ssl=1;AuthMech=3;httpPath=/sql/1.0/warehouses/624b8981d8e6e924;). The browser's address bar shows the URL: https://adb-2262725646607275.15.azuredatabricks.net/sql/warehouses/624b8981d8e6e924/connectionDetails?o=2262725646607275.

12) Dashboards are created for insights

